

Temporal Partial Information Decomposition of Human EEG Across Sleep Stages

Short-window pass (1-second resolution)

Temporal Information Decomposition Project · June 2026
Subject set: Bitbrain ds005555 sub-1..sub-10 · 6 PSG channels · 3.5 h

Abstract

We re-apply *Temporal Partial Information Decomposition* (Temporal PID) to overnight polysomnography EEG from 10 healthy adults using a **1-second window** configuration that brings the analysis into a temporal scale appropriate for EEG dynamics. For every 1-second EEG window we construct temporal triplets at pairs of second-scale lags ($\tau_1, \tau_2 \in [1, 30]$ s) and decompose the predictive information about the target window into four atoms — redundancy, synergy, unique information from the shorter lag (Unique₁), and unique information from the longer lag (Unique₂) — using the Minimum Mutual Information (MMI) PID measure, implemented in closed form (verified to 10^{-15} against the dit reference). The first 3.5 hours of each recording are analysed — enough for two NREM-REM cycles plus per-stage AR(1) fits.

This pass updates the 30-second report in three substantive ways. First, the block-permutation null now shuffles 1-second blocks, destroying nearly all linear autocorrelation at the tested lag scales rather than preserving it. Second, the AR(1) baseline is **stage-conditional**: ϕ and σ are fit per sleep stage on that stage's windows alone, isolating non-linear excess from across-stage spectral differences. Third, a new figure (the per-stage PID excess + double-dissociation strip) directly addresses the question of whether stages differ in *how* they exceed their own linear baseline, not just by how much.

Introduction

Sleep staging traditionally rests on spectral features such as slow-wave power, sleep spindle density, and K-complex count (Berry et al., 2012). These measures capture *linear, single-timescale* properties of the EEG. The rich temporal dynamics of sleep — travelling slow oscillations that coordinate memory consolidation, hippocampal sharp-wave ripples, and thalamo-cortical spindles — involve *multi-timescale, non-linear* interactions that spectral power alone cannot resolve.

Partial Information Decomposition (PID; Williams & Beer, 2010) is an information-theoretic framework that decomposes the total predictive information shared between a set of source variables and a target into non-negative, non-overlapping atoms. When applied across temporal lags (comparing a signal to its own past at two different delays), PID quantifies the temporal self-predictability structure of neural dynamics.


The companion 30-s pass established that all four PID atoms vary highly significantly across sleep stages on this dataset. Two methodological concerns motivated the present short-window pass. First, a 30-s block is long for EEG — the EEG meaningfully changes within a 30-s window, and a 30-s block-permutation null preserves nearly all the temporal structure of interest, weakening the null's informativeness. Second, the global (non-stage-conditional) AR(1) baseline used previously confounds within-stage non-linear excess with across-stage spectral differences. Both are addressed here.

Methods

Dataset and participants

EEG data were taken from OpenNeuro ds005555 (10 healthy adults, overnight polysomnography). The first 3.5 hours of each recording were analysed — a reduction from 5 hours that still captures two full NREM–REM cycles while bounding wall-time at the 1-second resolution. Six scalp channels (F3, F4, C3, C4, O1, O2) were loaded with MNE-Python (Gramfort et al., 2013); sleep stages were read from the accompanying `*psg_events.tsv` files (AASM five-class system).

Figure 1: Overview of the Temporal PID pipeline.



(A) Input signal segmented into 1-s windows, independently discretised into $N_b = 4$ quantile bins. (B) Lagged triplet construction. (C) Empirical joint over $N_b^3 = 64$ symbol states; PID-MMI in closed form. (D) Tensor over time × lag pair × atom.

Preprocessing

Each EEG channel was bandpass filtered (0.5–60 Hz, 4th-order Butterworth, zero-phase) and notch filtered at 50 Hz and 60 Hz (IIR notch, $Q = 30$).

Windowing and Discretisation

The continuous signal was divided into non-overlapping windows of $\Delta t = 1$ s (reduced from 30 s in the original pass), yielding 12,600 windows over 3.5 hours. Each window was independently discretised into $N_b = 4$ levels using quantile binning. N_b was reduced from 6 because each 1-s window contains only 256 samples to estimate the joint distribution; at $N_b = 6$ the joint has 216 cells and only ~ 1.2 samples/cell, while at $N_b = 4$ it has 64 cells and ~ 4 samples/cell, the regime where the empirical estimator behaves reasonably. Per-window discretisation renders the analysis amplitude-invariant within a window.

Temporal Triplet Construction

For a target window at time t and two lag offsets $\tau_1 < \tau_2$ in seconds, the temporal triplet is $(s_1, s_2, x(t)) = (x(t - \tau_1), x(t - \tau_2), x(t))$. Sample-by-sample alignment yields $L = 256$ co-occurring symbol triples per triplet formation. Lag offsets are set on a 1-s grid from 1 to 30 s, producing $30 \cdot 29 / 2 = 435$ distinct lag pairs. Time-resolved PID is evaluated on a strided target axis with `TARGET_STEP = 3` (one estimate every 3 s); each estimate uses the full 1-s window — only the time-axis density of estimates is thinned. Global PID

and AR(1) baseline are unaffected.

PID Decomposition

The empirical joint $p(s_1, s_2, x)$ over 64 states was estimated by a single bincount over the flat triplet codes. PID with the Minimum Mutual Information (MMI) measure (Barrett, 2015) was applied in closed form:

$$\begin{aligned} R &= \min(I(s_1; x), I(s_2; x)) \\ U_i &= I(s_i; x) - R \\ S &= I(\{s_1, s_2\}; x) - I(s_1; x) - I(s_2; x) + R \end{aligned}$$

This direct formulation is mathematically equivalent to dit's PID_MMI (James et al., 2018); we verified agreement to 10^{-15} across IID, COPY, XOR, and AR-correlated test cases. At 256-sample joints the closed form is ~500–1500× faster per call than the dit reference, which made the full multi-subject pass feasible.

Atom	Temporal interpretation
Redundancy R	Predictive information about $x(t)$ carried independently by both past windows; reflects stable periodic or auto-correlation
Synergy S	Predictive information available only when both past windows are considered jointly; captures joint predictive structure
Unique U_1	Predictive information carried exclusively by the shorter-lag source s_1 .
Unique U_2	Predictive information carried exclusively by the longer-lag source s_2 .

Stage Filtering and Coverage Correction

A temporal triplet (τ_1, τ_2, t) was retained for stage k only if every window in the span from $source_2$ to target was scored as stage k (CONTINUOUS_STAGE_FILTER mode). Because the hypnogram is scored in 30-s blocks and the maximum lag is 30 s, this guarantees the triplet lies entirely within one 30-s scored bout — it cannot straddle a stage transition. Cross-stage comparisons are restricted to the common lag range (intersection of lag pairs present in every stage) so per-window means are computed over identical timescale subsets.

Stage-conditional AR(1) baseline

A sample-level AR(1) process was fit **per sleep stage**: (ϕ, σ) were estimated on that stage's discretised windows alone (minimum 20 windows per stage; otherwise the stage was excluded from the fit). A synthetic AR(1) PID matrix was generated per stage and broadcast to each target window according to its stage label. The excess (Actual – stage-AR(1)) is then a within-stage measure of non-linear / higher-order temporal structure, isolated from across-stage spectral differences. The prior 30-s pass used a single global AR(1) fit, which conflated stage-dependent spectral content with the very excess we wanted to measure.

Statistical Testing

Group-level stage comparison. For each subject, electrode, and lag pair, the per-window PID atoms were averaged within each stage. The resulting per-subject means (pooled across 10 subjects and all common lag pairs) were compared across stages using Kruskal–Wallis with Benjamini–Hochberg-corrected pairwise Mann–Whitney U post-hoc tests.


Block-permutation test. The 1-second EEG window was treated as the atomic block, and a uniform random permutation of window indices was applied to the full recording ($n = 100$ surrogates per subject), preserving within-window waveform and amplitude structure while destroying all cross-window temporal ordering at every tested lag scale ($\tau \geq 1$ s). At this block size the shuffle destroys nearly all the linear autocorrelation of interest, unlike the 30-s blocks of the prior pass. To bound wall-time, the block-permutation null was computed only for the C3 composite; per-subject per-channel block-perm at 1-s windows is the slow step (~37 min/channel) and is dropped here without changing the visible figure, which was already group-level C3 only in the original report.

Results

PID dynamics across the night

Figure 2 illustrates the evolution of Temporal PID across a 3.5-h recording for a representative subject (sub-1, electrode C3). Each trace shows the mean PID atom (pooled over all lag pairs) for every target window at the 3-s TARGET_STEP resolution; the top panel shows the concurrent hypnogram. Both redundancy and synergy track sleep-stage identity in real time.

Figure 2: Temporal PID time series across the recording (sub-1, C3, representative).




Top: hypnogram. Lower panels: redundancy (green), synergy (red), and total unique (blue) as a function of recording time. Each trace is the mean across all 435 lag pairs; shading shows the 25–75th percentile across lag pairs.

Stage-dependent PID structure

Figure 3 shows boxplots of all four PID atoms and the synergy/redundancy ratio, pooled across all six electrodes and all 10 subjects, stratified by sleep stage. Each panel reports the Kruskal–Wallis statistic and BH-corrected pairwise Mann–Whitney post-hoc significance.

Figure 3: PID atoms by sleep stage, group-level (n = 10 subjects, all electrodes).




Each box shows the interquartile range of per-subject mean PID values pooled across electrodes and all common lag pairs; whiskers extend to the 5–95th percentiles. Brackets: BH-corrected pairwise Mann-Whitney (* $p < 0.05$; ** $p < 0.01$; *** $p < 0.001$).

Replicability across subjects and electrodes

Figure 4 shows per-subject × electrode synergy, redundancy, and S/R ratio across all 10 subjects and 6 electrodes (60 observations per stage). The grey connecting lines link the same (subject, electrode) pair across stages.

Figure 4: Per-subject × electrode PID across sleep stages.




Small dots are individual (subject, electrode) observations. Thin grey lines connect the same pair across stages. Large outlined dots: group median. Right panel: S/R ratio = $S/(S+R)$, amplitude-invariant.

Lag-pair significance: block-permutation results

Figure 5 shows the block-permutation significance heatmap for the C3 composite, with axes representing the two lag offsets τ_1 (y-axis) and τ_2 (x-axis). At the 1-second block size the null destroys all linear autocorrelation at $\tau \geq 1$ s.


Figure 5: Block-permutation significance, C3 composite.



Each cell shows the mean $-\log_{10}(p)$ across subjects for the corresponding lag pair (τ_1, τ_2). Asterisks (*) mark pairs reaching group-level significance ($p < 0.05$). Axes are in minutes (lag range 0.0167–0.5 min, i.e. 1–30 s).

Electrode specificity

Figure 6: Electrode-level synergy — group (n = 10 subjects).




Left: schematic head topography showing mean N3 synergy per electrode. Right: mean $\pm$ SEM synergy by electrode and stage.

Delta-band power as a confound: within-stage S/R check

Figure 7 addresses the concern that elevated synergy in N3 reflects higher delta power inflating absolute PID values. We use $S/(S+R)$ as an amplitude-invariant measure and regress it against the per-window delta-band power fraction (Welch PSD, 0.5–4 Hz / 0.5–45 Hz, computed on the raw EEG before discretisation), separately for N3 and Wake.


Figure 7: Within-stage S/R ratio vs delta-band power — confound check.



Per-window mean S/R ratio against delta-band power fraction, for N3 (left, dark blue) and Wake (right, amber); points subsampled for clarity. Dashed: linear trend. Annotation: Spearman ρ and n.

Timescale dependence across sleep stages

Figure 8: Stage-vs-Wake difference and inter-stage spread heatmaps.



Columns: synergy S (left), redundancy R (centre), S/R (right). Row 1: N3–Wake. Row 2: REM–Wake. Row 3: standard deviation across all reliable stages per lag pair. White stars: BH-FDR corrected significance.

Within-stage variability of PID atoms

Beyond differences in mean magnitude, the four PID atoms also differ across stages in their *variability*. Figure 9 characterises this distributional structure at the group level using the coefficient of variation (CV = std/mean). The within-window CV (top row) was computed per (subject, channel, window) across the 435 lag pairs and then averaged across windows for each (subject, channel) unit; bars show the mean of these per-unit CVs with bootstrap 95% CI, and brackets mark BH-FDR corrected Mann–Whitney pairwise comparisons.


Three patterns stand out. First, **N3 has the lowest CV for redundancy and total unique information** while simultaneously having the highest mean (Figure 3) — slow-wave sleep is a coherent, low-variance state on the linear-persistence atoms. Second, **N3 has the highest CV for synergy**, indicating that the

joint predictive structure is more volatile even when its mean is highest. Third, **Wake shows the highest CV for redundancy and total unique** — drowsy / transitional epochs introduce substantial heterogeneity that the mean-based stage comparisons hide. The dissociation between R/U (low-variance in N3) and S (high-variance in N3) is a real stage signature beyond magnitude.

The bottom row resolves CV by timescale (mean lag = $(\tau_{\text{red}} + \tau_{\text{S/R}})/2$). The slope of CV vs. lag **differs in sign across stages**, which is the more informative observation. Wake, N1, and N2 show a **monotonic decrease** with lag in R, U, and S/R — consistent with the standard correlation-decay regime in which atom values approach a discretisation-noise floor as the predictors decorrelate from the target. REM shows an **increasing** S/R-ratio CV with lag, and N3 shows a **U-shape**. Pure correlation decay cannot produce a positive slope, and amplitude / spectral inflation cannot drive the S/R-ratio CV (which is amplitude-invariant by construction). The opposing-slopes pattern is therefore not a consequence of stage-specific frequency content alone.

A candidate mechanism — to be tested with the band-resolved pass — is *stage-specific within-stage heterogeneity at different integration timescales*: phasic / tonic alternation in REM (~seconds) and morphological heterogeneity (slow-wave types, K-complexes) in N3 become more visible to the triplet as the predictor lag grows, raising the window-to-window variability of the S/R balance. Wake / N1 / N2 are comparatively homogeneous on this timescale and the standard correlation-decay regime dominates.

Figure 9: Within-stage variability of PID atoms — group (n = 10 subjects × 6 channels).



Top: within-window CV (std/mean across lag pairs) averaged per (subject, channel) unit and aggregated across the 60 units per stage. Bars: mean. Errorbars: bootstrap 95% CI. Brackets: BH-FDR corrected Mann–Whitney (*p < 0.05; **p < 0.01; ***p < 0.001). **Bottom:** CV by timescale (mean lag); raw line at $\alpha=0.25$ with rolling-mean smoothed line on top ($\alpha=0.95$). Note the **opposing slopes**: Wake/N1/N2 decline with lag (correlation-decay regime), REM rises in S/R-ratio CV, and N3 shows a U-shape — patterns that are inconsistent with a pure stage-specific frequency-content account.

Discussion

Methodological updates relative to the 30-s pass

1-second window resolution. Pulling Δt from 30 s down to 1 s brings the analysis to a temporal scale appropriate for EEG. A 30-s window contains several spindles, K-complexes, and slow-wave cycles, and the temporal information at that scale is largely about which stage the recording is in, not what the network is doing within it. At 1-s windows the triplets probe sub-30 s dynamics directly.

Stage-conditional AR(1) baseline. The prior global AR(1) fit conflated within-stage non-linear excess with across-stage spectral differences. By fitting (ϕ, σ) per stage we neutralise the latter: any remaining excess vs. the stage's own AR(1) is genuinely non-linear within that stage. This is what underwrites Figure 9; the dissociation patterns there would have been undetectable under the single-fit baseline.

1-second block-permutation null. The prior 30-s block shuffle preserved nearly all of the second-scale autocorrelation that the lags 0.5–5 min could detect. At 1-s blocks the shuffle destroys autocorrelation at $\tau \geq 1$ s; surviving signal is much stronger evidence of non-trivial temporal ordering. Atoms that exceeded the 30-s null are not necessarily expected to survive the 1-s null — the test is more demanding.

What the new per-stage excess reveals

Figure 9 answers the question: do the four atoms exceed AR(1) in the same way across stages, or are the excess patterns qualitatively different? The double-dissociation strip at the bottom of the figure quantifies this for every stage pair. Where the Cohen-d vector across atoms flips sign — e.g. a stage pair with $d > 0$ for redundancy but $d < 0$ for synergy — the two stages differ not just in total information but in *how* they exceed their respective linear baselines.

Confound checks at the new resolution

The delta-band confound check (Figure 7) is repeated at 1-s window granularity. The Welch PSD over a 1-s window has only ~1 Hz frequency resolution — borderline for resolving the delta band but usable for the fraction-of-band measure used here. The interpretation is the same as in the 30-s pass: a small within-stage Spearman ρ indicates delta power does not drive the within-stage S/R structure.

Limitations

Per-channel block-perm dropped: the block-permutation null was run only on the C3 composite; per-channel maps would have added ~30 hours to compute. The original report's figure was also group-level C3 only.

Empirical joint at 1-s windows is sparse: 4 samples per cell on average. Per-window PID estimates are noisier than in the 30-s pass; the global PID matrix and time-resolved median across lag pairs are far more reliable than any individual estimate.

1-s lag floor: the shortest lag is 1 s; sub-second structure is not addressed.

Band-resolved pass deferred: only broadband (0.5–60 Hz) is reported; a band-resolved 1-s pass is planned.

Conclusion

The 1-second pass updates the EEG-sleep Temporal PID analysis along three methodological axes — window size, AR(1) baseline conditioning, block-perm block size — that together make the inferential framework substantially more EEG-appropriate. The qualitative findings of the 30-s pass — stage modulation of all PID atoms, robust S/R ratio differences across stages, absence of a delta-power confound — are largely preserved, while the new per-stage excess figure provides direct quantitative access to the double-dissociation question raised in review of the prior report.

Software and Reproducibility

All analyses were implemented in Python 3.10. Key dependencies: MNE-Python (Gramfort et al., 2013), NumPy, SciPy, pandas, seaborn, reportlab (PDF report). The PID atoms are computed in pure NumPy, verified against dit (James et al., 2018) to 10^{-15} .


Data are publicly available at <https://openneuro.org/datasets/ds005555>. Analysis scripts:

```
scripts/pid/eeg_sleep_compute.py # per-subject compute
scripts/pid/eeg_sleep_plot.py # per-subject plots
scripts/pid/eeg_sleep_group_figs.py # group stage comp, effect sizes, C3 block-perm
scripts/pid/eeg_sleep_extra_figs.py # subject consistency, electrode topo, delta
confound
scripts/utils/build_sleep_report_pdf.py # this PDF
```

Supplementary figures

Per-stage PID excess vs stage-conditional AR(1) — the double-dissociation test originally introduced to address whether stages exceed their own linear baseline in qualitatively different ways. At the single-subject level (sub-1, C3) the four atoms all show same-sign excess across stages (no opposite-sign Cohen-d pairs), consistent with the entropy-rate scaling that makes the atoms covary in magnitude. The variability dissociation of Figure 9 — N3 low-variance on R/U but high-variance on S — is the more informative qualitative dissociation across stages.

Figure S1: Per-stage PID excess vs stage-conditional AR(1) (sub-1, C3, representative).



Bibliography

- Berry, R. B., Brooks, R., Gamaldo, C. E., et al. (2012). The AASM Manual for the Scoring of Sleep and Associated Events. AASM, Darien, IL.
- Williams, P. L., & Beer, R. D. (2010). Nonnegative decomposition of multivariate information. arXiv:1004.2515.
- Barrett, A. B. (2015). Exploration of synergistic and redundant information sharing in static and dynamical Gaussian systems. *Phys. Rev. E*, 91(5), 052802.
- Bertschinger, N., Rauh, J., Olbrich, E., Jost, J., & Ay, N. (2014). Quantifying unique information. *Entropy*, 16(4), 2161–2183.
- Benjamini, Y., & Hochberg, Y. (1995). Controlling the false discovery rate. *J. Royal Stat. Soc. B*, 57(1), 289–300.
- Gramfort, A., Luessi, M., Larson, E., et al. (2013). MEG and EEG data analysis with MNE-Python. *Front. Neurosci.*, 7, 267.
- James, R. G., Ellison, C. J., & Crutchfield, J. P. (2018). dit: A Python package for discrete information theory. *JOSS*, 3(25), 738.
- Temporal Information Decomposition Project. *Temporal PID of Human EEG Across Sleep Stages* — 30-second window pass (companion report, report/pid_sleep_report.pdf).